

Proceedings

**14th International Conference on
Data Engineering**

Proceedings

14th International Conference on Data Engineering

February 23-27, 1998

Orlando, Florida

Sponsored by

IEEE Computer Society Technical Committee on
Data Engineering



Los Alamitos, California

Washington

•

Brussels

•

Tokyo

Copyright © 1998 by The Institute of Electrical and Electronics Engineers, Inc.
All rights reserved

Copyright and Reprint Permissions: Abstracting is permitted with credit to the source. Libraries may photocopy beyond the limits of US copyright law, for private use of patrons, those articles in this volume that carry a code at the bottom of the first page, provided that the per-copy fee indicated in the code is paid through the Copyright Clearance Center, 222 Rosewood Drive, Danvers, MA 01923.

Other copying, reprint, or republication requests should be addressed to: IEEE Copyrights Manager, IEEE Service Center, 445 Hoes Lane, P.O. Box 133, Piscataway, NJ 08855-1331.

The papers in this book comprise the proceedings of the meeting mentioned on the cover and title page. They reflect the authors' opinions and, in the interests of timely dissemination, are published as presented and without change. Their inclusion in this publication does not necessarily constitute endorsement by the editors, the IEEE Computer Society, or the Institute of Electrical and Electronics Engineers, Inc.

IEEE Catalog Number 98CB36164
ISBN: 0-8186-8289-2 (Softbound)
ISBN: 0-7803-4750-1 (Casebound)
ISBN: 0-8186-8291-4 (Microfiche)
ISSN: 1063-6382
PR8289

Additional copies may be ordered from:

IEEE Computer Society
Customer Service Center
10662 Los Vaqueros Circle
P.O. Box 3014
Los Alamitos, CA 90720-1314
Tel: + 1-714-821-8380
Fax: + 1-714-821-4641
E-mail: cs.books@computer.org

IEEE Service Center
445 Hoes Lane
P.O. Box 1331
Piscataway, NJ 08855-1331
Tel: + 1-908-981-1393
Fax: + 1-908-981-9667
mis.custserv@computer.org

IEEE Computer Society
13, Avenue de l'Aquilon
B-1200 Brussels
BELGIUM
Tel: + 32-2-770-2198
Fax: + 32-2-770-8505
euro.ofc@computer.org

IEEE Computer Society
Ooshima Building
2-19-1 Minami-Aoyama
Minato-ku, Tokyo 107
JAPAN
Tel: + 81-3-3408-3118
Fax: + 81-3-3408-3553
tokyo.ofc@computer.org

Editorial production by Regina Spencer Sipple

Cover art production by Joseph Daigle/Studio Productions

Printed in the United States of America by The Printing House



Table of Contents

14th International Conference on Data Engineering (ICDE '98)

General Chair's Message	xii
Program Co-Chairs' Message	xiii
Organizing Committee	xv
Steering Committee	xvii
Program Committee Vice Chairs	xviii
Program Committee	xix
External Reviewers	xxi

Keynote Address

Databases: The Next Generation

Andrew Mendelsohn, Oracle Corporation

Session 1: Semistructured Data

Chair: Boris Chidlovskii, Rank Xerox Research Center

Representing and Querying Changes in Semistructured Data	4
<i>S.S. Chawathe, S. Abiteboul, and J. Widom</i>	
Optimizing Regular Path Expressions Using Graph Schemas	14
<i>M. Fernández and D. Suciu</i>	
WebOQL: Restructuring Documents, Databases and Webs	24
<i>G.O. Arocena and A.O. Mendelzon</i>	

Session 2: Mobile and Disconnected Database Operations

Chair: Anupam Joshi, Missouri University

Concurrent Operations in a Distributed and Mobile Collaborative Environment	36
<i>M. Suleiman, M. Cart, and J. Ferrié</i>	
Grouping Techniques for Update Propagation in Intermittently Connected Databases	46
<i>S. Mahajan, M. Donahoo, S. Navathe, M. Ammar, and S. Malik</i>	
Cache Management for Mobile Databases: Design and Evaluation	54
<i>B.Y. Chan, A. Si, and H.V. Leong</i>	

Industry Session 1: Database Engines

Chair: Surajit Chaudhuri, Microsoft Research

Messaging/Queuing in Oracle8	66
<i>D. Gawlick</i>	
The New Database Imperatives	69
<i>G. Graefe</i>	

Panel Session 1:**WWW and the Internet — Did We (the Database Community) Miss the Boat?..... 74****Chair: Michael Rabinovich, AT&T Laboratories****Panelists: Mic Bowman, Transarc***Alon Levy, University of Washington**Susan Malaika, IBM Santa Teresa Lab**Alberto Mendelson, University of Toronto**Hector Garcia-Molina, Stanford University***Session 3: Parallel Database Systems****Chair: Kia Makki, University of Southwestern Louisiana****Performance Analysis of Parallel Hash Join Algorithms on a Distributed Shared
Memory Machine: Implementation and Evaluation on HP Exemplar SPP 1600 76***M. Nakano, H. Imai, and M. Kitsuregawa***Parallelizing Loops in Database Programming Languages..... 86***D.F. Lieuwen***Cyclic Allocation of Two-Dimensional Data..... 94***S. Prabhakar, K. Abdel-Ghaffar,**D. Agrawal, and A. El Abbadi***Industry Session 2: Data Mining****Chair: Meichun Hsu, Hewlett Packard Labs****Employing Intelligent Agents for Knowledge Discovery 104***E. Stahl***Industry Applications of Data Mining: Challenges & Opportunities..... 105***E. Simoudis***Session 4: Queries and Views****Chair: Ed Fernandez, Florida Atlantic University****Fuzzy Triggers: Incorporating Imprecise Reasoning into Active Databases 108***A. Wolski and T. Bouaziz***Graph Structured Views and Their Incremental Maintenance..... 116***Y. Zhuge and H. Garcia-Molina***Query Folding with Inclusion Dependencies 126***J. Gryz***On Querying Spreadsheets 134***L.V.S. Lakshmanan, S.N. Subramanian,**N. Goyal, and R. Krishnamurthy***Session 5: Recovery and Buffer Management****Chair: Wojtek Cellary, University of Economics at Poznan****Data Logging: A Method for Efficient Data Updates in Constantly
Active RAID's..... 144***E. Gabber and H.F. Korth***Persistent Applications Using Generalized Redo Recovery..... 154***D.B. Lomet*

The Effect of Buffering on the Performance of R-Trees	164
<i>S.T. Leutenegger and M.A. Lopez</i>	

Industry Session 3: OLAP and Data Warehousing

Chair: Kyuseok Shim, Bell Laboratories

Red Brick Vista™: Aggregate Computation and Management	174
<i>L.S. Colby, R.L. Cole, E. Haslam, N. Jazayeri, G. Johnson, W.J. McKenna, L. Schumacher, and D. Wilhite</i>	
Outstanding Challenges in OLAP	178
<i>J.A. Bedell</i>	
Ending the ROLAP/MOLAP Debate: Usage Based Aggregation and Flexible HOLAP	180
<i>C. Salka</i>	

Keynote Address

Safeguarding and Charging for Information on the Internet	182
<i>Hector Garcia-Molina, Steven P. Ketchpel, and Narayanan Shivakumar, Stanford University</i>	

Session 6: Spatial and Temporal Databases

Chair: Mostafa Bassiouni, University of Central Florida

Point- Versus Interval-Based Temporal Data Models	192
<i>M.H. Böhlen, R. Busatto, and C.S. Jensen</i>	
Efficient Retrieval of Similar Time Sequences under Time Warping	201
<i>B.-K. Yi, H.V. Jagadish, and C. Faloutsos</i>	
Fast Nearest Neighbor Search in High-Dimensional Space	209
<i>S. Berchtold, B. Ertl, D.A. Keim, H.-P. Kriegel, and T. Seidl</i>	

Session 7: Indexing in Data Warehousing

Chair: Ravi Mukkamala, Old Dominion University

Encoded Bitmap Indexing for Data Warehouses	220
<i>M.-C. Wu and A.P. Buchmann</i>	
Coarse Indices for a Tape-Based Data Warehouse	231
<i>T. Johnson</i>	
Array-Based Evaluation of Multi-Dimensional Queries in Object- Relational Database Systems	241
<i>Y. Zhao, K. Ramasamy, K. Tufte, and J.F. Naughton</i>	

Industry Session 4: Multimedia and Workflow Systems

Chair: Hector Garcia-Molina, Stanford University

Content-Based Multimedia Information Management	252
<i>R. Jain</i>	
Building a Robust Workflow Management System with Persistent Queues and Stored Procedures	254
<i>F. Leymann and D. Roller</i>	

An Extended Object-Oriented Database Approach to Networked Multimedia Applications	259
<i>H. Ishikawa, K. Kato, M. Ono, N. Yoshizawa, K. Kubota, and A. Kanaya</i>	

Session 8: Video and Image Databases

Chair: Ming-Syan Chen, National Taiwan University	
Distributed Video Presentations	268
<i>E. Hwang, V.S. Subrahmanian, and B. Prabhakaran</i>	
Query Processing in a Video Retrieval System	276
<i>K.L. Liu, P. Sistla, C. Yu, and N. Rishe</i>	
SEMOG: A Hybrid Object-based Image Database System and Its Modeling, Language, and Query	284
<i>W.-S. Li and K.S. Candan</i>	

Panel Session 2:

Data Warehousing Lessons from Experience	294
Chair: Patrick O'Neil, University of Massachusetts, Boston	
Panelists: <i>Richard Winter, Winter Consulting</i>	
<i>Clark French, Sybase IQ</i>	
<i>Dan Crowley, Informix Software</i>	
<i>William McKenna, Red Brick Systems</i>	

Industry Session 5: The Web and Databases

Chair: Alex Buchmann, T.H. Darmstadt	
Junglee: Integrating Data of All Shapes and Sizes	296
<i>A. Gupta</i>	
Virtual Database Technology	297
<i>A. Gupta, V. Harinarayan, and A. Rajaraman</i>	
Data-Intensive Intra- & Internet Applications — Experiences Using Java and CORBA in the World Wide Web	302
<i>J. Sellentin and B. Mitschang</i>	

Panel Session 3:

Migrating Legacy Databases and Applications	314
Chair: Bhavani Thiraisingham, MITRE Corporation	
Panelists: <i>Sandra Heiler, GTE Laboratories</i>	
<i>Arnon Rosenthal, MITRE Corporation</i>	
<i>Susan Malaika, IBM Santa Teresa Lab</i>	
<i>Others TBD</i>	

Session 9: Architectural Issues in Data Mining Systems

Chair: Lois Delcambre, Oregon Graduate Center	
A Tightly-Coupled Architecture for Data Mining	316
<i>R. Meo, G. Psaila, and S. Ceri</i>	

A Distribution-Based Clustering Algorithm for Mining in Large Spatial Databases	324
<i>X. Xu, M. Ester, H.-P. Kriegel, and J. Sander</i>	

Session 10: Cooperation and Workflow Management

<i>Chair: Niki Pissinou, University of Southwestern Louisiana</i>	
Failure Handling and Coordinated Execution of Concurrent Workflows	334
<i>M. Kamath and K. Ramamritham</i>	
Methodical Restructuring of Complex Workflow Activities	342
<i>L. Liu and C. Pu</i>	
Leveraging Mediator Cost Models with Heterogeneous Data Sources	351
<i>H. Naacke, G. Gardarin, and A. Tomasic</i>	

Session 11: Access Structures

<i>Chair: Thomas Mueck, University of Wien</i>	
The LSD ^h -tree: An Access Structure for Feature Vectors	362
<i>A. Henrich</i>	
Compressing Relations and Indexes	370
<i>J. Goldstein, R. Ramakrishnan, and U. Shaft</i>	
Generalizing "Search" in Generalized Search Trees	380
<i>P.M. Aoki</i>	

Session 12: Association Rules and Dependencies

<i>Chair: X. Hu, Sybase Inc.</i>	
Efficient Discovery of Functional and Approximate Dependencies Using Partitions	392
<i>Y. Huhtala, J. Kärkkäinen, P. Porkka, and H. Toivonen</i>	
Online Generation of Association Rules	402
<i>C.C. Aggarwal and P.S. Yu</i>	
Cyclic Association Rules	412
<i>B. Özden, S. Ramaswamy, and A. Silberschatz</i>	

Session 13: Concurrency Control

<i>Chair: Vijay Atluri, Rutgers University</i>	
Asynchronous Version Advancement in a Distributed Three Version Database	424
<i>H.V. Jagadish, I.S. Mumick, and M. Rabinovich</i>	
Design and Performance of an Assertional Concurrency Control System	436
<i>A.J. Bernstein, D.S. Gerstl, W.-H. Leung, and P.M. Lewis</i>	
Dynamic Granular Locking Approach to Phantom Protection in R-trees	446
<i>K. Chakrabarti and S. Mehrotra</i>	

Plenary Panel Session:

Future Directions in Database Research.....	456
--	------------

Chair: Surajit Chaudhuri, Microsoft Corporation

Panelists: Hector Garcia-Molina, Stanford University

Hank Korth, Bell Laboratories

Guy Lohman, IBM Almaden Research Center

David Lomet, Microsoft Research

David Maier, Oregon Graduate Institute

Session 14: Query Processing in Spatial Databases

Chair: Sharma Chakravarthy, University of Florida

Processing Incremental Multidimensional Range Queries in a Direct Manipulation Visual Query Environment	458
--	-----

S. Hibino and E. Rundensteiner

High Dimensional Similarity Joins: Algorithms and Performance Evaluation	466
--	-----

N. Koudas and K.C. Sevcik

Cost Models for Join Queries in Spatial Databases	476
---	-----

Y. Theodoridis, E. Stefanakis, and T. Sellis

Session 15: Mining of Association Rules

Chair: Michael Soo, University of South Florida

Mining Association Rules: Anti-Skew Algorithms	486
--	-----

J.-L. Lin and M.H. Dunham

Mining for Strong Negative Associations in a Large Database of Customer Transactions	494
---	-----

A. Savasere, E. Omiecinski, and S. Navathe

Mining Optimized Association Rules with Categorical and Numeric Attributes.....	503
--	-----

R. Rastogi and K. Shim

Session 16: Server-Based Database Systems

Chair: Anoop Singhal, AT&T Laboratories

Remote Load-Sensitive Caching for Multi-Server Database Systems	514
---	-----

S. Venkataraman, J.F. Naughton, and M. Livny

DB-MAN: A Distributed Database System Based on Database Migration in ATM Networks	522
--	-----

*T. Hara, K. Harumoto, M. Tsukamoto,
and S. Nishio*

Network Latency Optimizations in Distributed Database Systems	532
---	-----

S. Banerjee and P.K. Chrysanthis

Session 17: Visualization of Multimedia Data

Chair: Tiziana Catarci, Università di Roma "La Sapienza"

Global Integration of Visual Databases	542
--	-----

*W. Chang, D. Murthy, A. Zhang,
and T.F. Syeda-Mahmood*

The Alps at Your Fingertips: Virtual Reality and Geoinformation Systems	550
<i>R. Pajarola, T. Ohler, P. Stucki,</i>	
<i>K. Szabo, and P. Widmayer</i>	
Design and Implementation of Display Specifications for Multimedia Answers	558
<i>C. Baral, G. Gonzalez, and T.C. Son</i>	

Session 18: Management of Objects

<i>Chair: Arbee Chen, National Tsing Hua University</i>	
Flattening an Object Algebra to Provide Performance	568
<i>P. Boncz, A.N. Wilschut, and M.L. Kersten</i>	
Back to the Future: Dynamic Hierarchical Clustering	578
<i>C. Zou, B. Salzberg, and R. Ladin</i>	
Cost and Imprecision in Modeling the Position of Moving Objects	588
<i>O. Wolfson, S. Chamberlain, S. Dao,</i>	
<i>L. Jiang, and G. Mendez</i>	

Research Prototype Demonstrations

ROL: A Prototype for Deductive and Object-Oriented Databases	598
<i>M. Liu, W. Yu, M. Guo, and R. Shan</i>	
A Graphical Editor for the Conceptual Design of Business Rules	599
<i>P. Lang, W. Obermair, W. Kraus, and T. Thalhammer</i>	
The Active HYpermedia Delivery System (AHYDS) using the PHASME Application-Oriented DBMS	600
<i>F. Andres and K. Ono</i>	
ECA Rule Support for Distributed Heterogeneous Environments	601
<i>S. Chakravarthy and R. Le</i>	
ZEBRA Image Access System	602
<i>S. Mudumbai, K. Shah, A. Sheth,</i>	
<i>K. Parasuraman, and C. Bertram</i>	

Author Index	603
---------------------------	------------